

# CQS Product Requirements

## CQS Product Requirements

### Overview

This document specifies the high-level requirements for CQS that all stakeholders can agree upon. One goal of this page is provide criterion for evaluating various design options.

### Functional Requirements

These assumes normal operating circumstances.

- Implement all actions in the Amazon SQS spec
- The Queues do not have to support FIFO behavior under any condition. No effort is done to make CQS behave like a traditional Queue data-type
- HTTP based protocol (as specified by Amazon SQS)
- Guaranteed delivery of all messages at least once
- Redelivery of messages that are not deleted in a configurable timeout.
- Redelivery of deleted messages during or after some failure conditions like Cassandra node failure.
- For the ReceiveMessage() API call, don't return the same message multiple times in more than 5% of the cases.
- A ReceiveMessage request for n messages might return less than n messages even though system has n messages in less than 10% of cases
- Message size no more than 64k

### Response Time Requirements

These assumes normal operating circumstances.

- For the Send/ReceiveMessage of 1 message, the response time should be ~10ms in 95% of cases within a data-center.
- For the Send/ReceiveMessage of 1 message, the response time should be ~100ms in 99% of the cases within a data-center
- For the Send/ReceiveMessage of 1 message, the response time should be ~8ms in the median cases
- The time between a succesfull response to SendMessage() and the time a RecieveMessage() can return that message must not be more than 500ms in the 95%

### Availability Requirements

All the below requirements assume that all the underlying infrastructure behaves like we expect in a failover scenario (Cassandra replication, load-balancers switching, network, etc)

- System needs to be highly available (99.999%) and deployed in multiple data centers
- No message loss in the event of node or data-center failure.
- In data-center failure case, The messages should be made available in a different data-center instantly with possible loss of hidden state.
- When a single-node failure occurs within a data-center ALL the messages should be made available in less than 60 seconds

### Usage Pattern

- In Normal mode of operation, message producers and message consumers of a queue are co-located in the same data-center.
- In data-center failure case, the messages need to be available in some other data-center where the consumers can be directed to consume the messages. The consumers would be redirected through a load-balancer.

## **Scalability Requirements**

- Unlimited number of queues
- Horizontally scalable. Adding more API and Cassandra nodes should increase throughput and capacity as linearly as feasible
- Must handle a peak load of 500k sendMessage() requests/sec
- Must have a peak load of 100k receiveMessage() requests/sec.

## **Security Requirements**

For the first version of SQS we are not planning on implementing the AddPermission command or the CreateQueue's 'Policy' attribute. In subsequent releases, we can do these.

## **Client Usage Recommendation**

Since the SQS will return messages out of order and could occasionally return messages previously returned by the Queue, the clients should be prepared to re-order messages and ignore duplicates.